



Overview

Board games present a unique challenge for AI agents. Coup, one of the most popular board games, is more challenging than others because deception plays a key role in gameplay. It is a multi-agent game and each agent would pick the optimal policy against other players with specific strategies. We modeled it as a Markov Decision Process (MDP) with unknown transition functions and trained an intelligent agent using Q-learning and feature extractions. The biggest challenge is the large state space.

Due to the complexity of the game, we re-range the game scope by introducing three simplifications:

1. Disable Ambassador's exchange action;
2. Only keep two cards for each character instead of three;
3. Only 3 players are enrolled in this game while the original version allows 2-6 players.

play			
Take one Action (If 10+ coins must choose to launch Coup)			
Character	Action	Effect	Counteraction
—	Income	Take 1 coin	×
—	Foreign Aid	Take 2 coins	×
—	Coup	Pay 7 coins Choose player to lose influence	×
Duke	Tax	Take 3 coins	Blocks Foreign Aid
Assassin	Assassinate	Pay 3 coins Choose player to lose influence	×
Ambassador	Exchange	Exchange cards with Court Deck	Blocks stealing
Captain	Steal	Take 2 coins from another player	Blocks stealing
Contessa	×	×	Blocks assassination

Models

[States]

Game state:

$s_{game} = [[["duke", 1], ["ambassador", 1]],$
 $["assassin", 1], ["contessa", 2]],$
 $["captain", 1], ["ambassador", 1]],$
 $[4,5,3], ["steal", 2, 1]$

Agent state:

$s_{a0} = [[["duke", 1], ["ambassador", 1]],$
 $1,$
 $2],$
 $[4,5,3], ["steal", 2, 1]$

[Actions]

$s_{a2} = [2, 1, ["captain", 1], ["ambassador", 1]], [4,5,3], None$


$Actions(s_{a2}) = \{["income", 2, 2], ["foreign aid", 2, 2],$
 $["assassinate", 2, 0], ["assassinate", 2, 1],$
 $["tax", 2, 2], ["steal", 2, 0], ["steal", 2, 1]\}$

$s_{a0} = [[["duke", 1], ["ambassador", 1]], 1, 2], [4,5,3], ["steal", 2, 1]$

$Actions(s_{a0}) = \{["doubt", 0, 2], ["no doubt", 0, 2]\}$

$Actions(s_{a1}) = \{["block steal", 1, 2], ["no block steal", 1, 2]\}$

[Transition function]



[Reward]

- Win the game → + 1000
- Opponents lose cards → + 100 /card
- Lose cards → - 500/ card

[Evaluation function]

$$\text{WinRate} = \frac{\# \text{win}}{\# \text{total trials}}$$

Algorithms

• Q-learning



1. Let the current state be s .
2. Select an action a to perform.
3. Let the reward received for performing a be r , and the resulting state be t .
4. Update $Q(s, a)$ to reflect the observation $\langle s, a, r, t \rangle$ as follows:

$$Q(s, a) = (1 - \alpha)Q(s, a) + \alpha(r + \gamma \max_{a'} Q(t, a'))$$
 where α is the current learning rate.
5. Go to step 1.




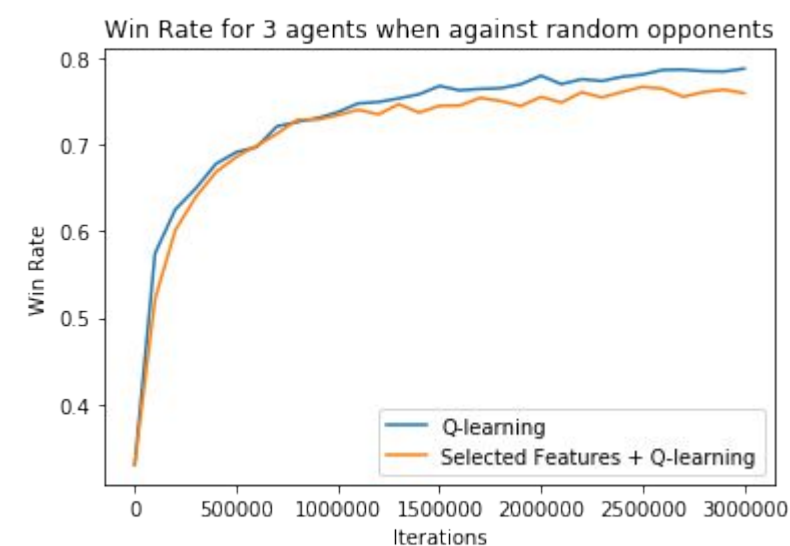

• Feature Extractions

- Convert complicated states to simpler selected features
- [tax, steal, assassinate, ... block assassinate, living cards of agents, coins left of agents]
- [1 0 ... 0 0 | 2 1 1 | 3 2 0]
 Actions Living Cards Coins left

• Epsilon Greedy

- fixed: $\epsilon = 0.2$

Experiments & Results



Experiments & Results

Iterations	Training	Testing
2 Agents	150,000	10,000
3 Agents	3,000,000	10,000

Runtime(s)	Q-learning	Features + QL
2 Agents	79.9	68.1
3 Agents	2350.9	2245.9

Simple Strategy vs. Simple Strategy

Conclusions

- Both our Q-learning algorithms provide substantial results. Trained agents perform impressively and reach high win rates.
- Q-learning with feature extractions is generally faster than Q-learning in the same experiment setting, but there is a tradeoff between win rate and running time.
- At a high level, actions distributions on states in learned policies are similar.
- The win rate of random vs. random is higher than simple vs. simple for both 2 and 3 agents.

Reference

[WD92] Christopher JCH Watkins and Peter Dayan. "Q-learning". In: Machine learning 8.3-4. (1992), pp. 279–292.
 [RN94] Gavin A Rummery and Mahesan Niranjan. Online Q-learning using connectionist systems. Vol. 37. University of Cambridge, Department of Engineering Cambridge, England, 1994.
 [Mni+15] Volodymyr Mnih et al. "Human-level control through deep reinforcement learning". In: Nature 518.7540 (2015), p. 529.
 [Sil+16] David Silver et al. "Mastering the game of Go with deep neural networks and tree search". In: nature 529.7587 (2016), p. 484.