



# Comparing Reinforcement Learning Techniques on carRacing-v0 Environment

Julius Stener & Victoria Valverde  
Stanford University



## Motivation

- The combination of Deep Learning and Reinforcement Learning have been successful at solving complicated riddles and extracting non-linear features from high dimensional media such as images<sup>1</sup>
- The novelty of the field and our interest in deep learning and reinforcement learning drove us to choose this project

## Problem Definition

- Compare how three continuous action space RL algorithms -- **PPO** (Proximal Policy Optimization), **DDPG** (Deep Deterministic Policy Gradient), and **TD3** (Twin Delayed DDPG) - perform the task of teaching a race car to drive around a track
- Tune hyperparameters to maximize performance

## Approach: Algorithm Description

**PPO**<sup>3</sup>: on-policy, e-greedy and Gaussian noise exploration method<sup>3</sup>


- Pros: uses SGD to optimize parameters, parallelizable<sup>3</sup>
- Cons: does not account for how exploration should be performed, L can still change significantly between updates

**DDPG**<sup>3</sup>: off-policy, Ornstein–Uhlenbeck process exploration method<sup>3</sup>

- Pros: Q-function used to learn the value of actions (not states)
- Cons: policy is only as good as the Q-function is accurate, learned Q-function commonly overestimates Q-values<sup>4</sup>

**TD3**<sup>4</sup>: improved version of DDPG

- Pros: improved through clipped double learning, delayed policy updates and target policy smoothing
- Cons: very compute intensive



Illustrative state processing from pixel image to linear vector with CNN


## Approach: Model Description

### CarRacing-v0 environment<sup>2</sup>:


- Game engine from OpenAI gym, part of the continuous control tasks in the Box2D simulator environments

We used a Markov Decision Process (**MDP**) to model our task:

- **Observation Space**: set of 96x96x3 vectors encoding the image on the right
- **Action Space**: set of 3D vectors of the form [s,t,b] where each letter specifies the steering angle (s), throttle (t) and brake (b)
- **Reward**: -0.1 for every frame and +1000/N for every track tile visited, where N is the total number of tiles in the track



## Results: Algorithm Performance & Noise



### PPO

- Trained E=1000 episodes, N=1000 parallel actors each collecting a default of T=8 timesteps of data. Each episode constructs surrogate loss on the NT timesteps of data and optimize using SGD
- Showed **considerable improvement** after 700 episodes, increasing at an average of 0.3 points per episode during the last 300 iterations
- Considerable amount of **noise** (specially at start) due to the nature of PPO and L being able to change significantly between updates

## Results: Algorithm Performance & Noise



### DDPG & TD3

- Trained for E=1000 episodes, T=1000 timesteps each choosing actions with exploration noise of 0.5, and batch size of N=100 from the replay buffer to update the policy
- Implementation was unable to learn, hence score stayed around -17.9 but we can still observe a considerable amount of noise

### Error Analysis

- Unable to make the race car move due to an unintended mismatch in the dimensionality of our NN's output and that of the action space
- Suspect needs tuning of hyperparameters as well as in more iterations to train in order to see a significant improvement

## Challenges & Future Directions

- **Further tuning parameters**: would potentially improve performance
- **Running on GPU**: would allow us to further tune and run tests more efficiently, which was the main challenge in our project
- **Discretizing action space**: would allow us to make the search and optimization faster/simpler
- **Removing the bottom panel**: (containing information such as lateral g-force, reward indicator and velocity) does not add meaningful information about current state. Decreasing the number pixels to convolve through has the potential to increase accuracy and speed
- **Setting a maximum on the rewards**: would stop the model from incentivizing high speeds that make car loose control in tight curves