



# Minecraft AI: Finding Water Sources

Lydia Chan<sup>1</sup> and Russell Tran<sup>1</sup>, {lchan528, tranrl}@stanford.edu

<sup>1</sup>CS221 Artificial Intelligence: Principles and Techniques, Stanford University

Stanford  
Computer Science

## Overview

**Problem:** Complete a search task in an unknown environment

**Motivation:** Autonomous robots for dangerous search and rescue missions

**Approach:** Find a source of water in a virtual environment in **Minecraft**



Fig 1. Waterfall in Minecraft



Fig 2. Large Minecraft Lake

## Dataset

### Dataset Sources

- **MALMO** to design a variety of interactive Minecraft worlds
- **MineRL** to interface with the environments using OpenAI Gym

### Dataset Environments

- **Basic:** Rectangular box made of a uniform material (stone)




Fig 3. Examples of a Basic Environment

- **Sparse:** flat, grassy terrain




Fig 4. Examples of a Sparse Environment

- **Dense:** hilly, jungle-based terrain with animals, deserts, or tundras



Fig 5. Examples of a Dense Environment

## Baseline

### Brute-force policy

- Agent ignores all sensory input
- systematically traverses every tile in the environment until it inevitably touches the target
- High standard but fails outside of basic environment




Fig 6. Illustrated Approach of Baseline

## Q-learning

### Naive approach

- Identity feature extractor for sensory (pixel) input
- Constant exploration rate of 30%

### Heuristics-based approach

- Feature extractor captures the value of blue pixel from each pixel in the 64x64 frame

## Deep Q-learning (DQN)

### Comparison with Q-learning

- Feature extractor was time-expensive
- We wanted to evolve the exploration rate
- Feature extractor (sensory input) is based on deep neural networks
- Linear exploration schedule transitions from a randomness exploration rate of 100% to 2% in more than 100,000 time steps
- Stochastic gradient descent was used to train the DQN

## Results

\*Reward: +6000 for touching water, -1 per step taken




Fig 7. Percentage of Time Spent Exploring per Episode




Fig 8. Mean Episode Reward

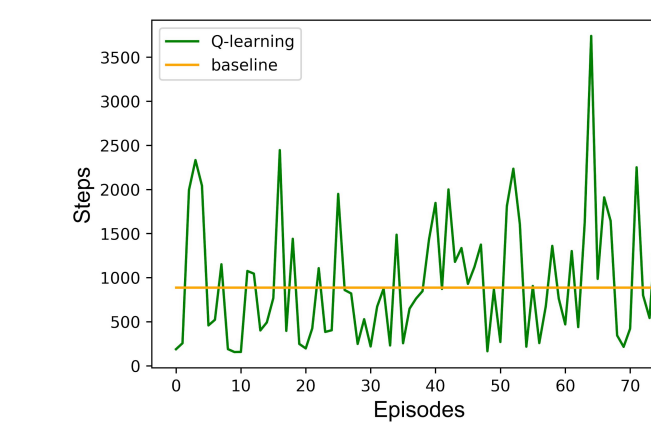



Fig 9. Number of Steps per Episode

## Discussion

### Memory Constraints

- Naive approach to Q-learning exceeds realistic memory constraints; better features are necessary

### Number of Steps per Episode

- No evident correlation: we may reconsider this as a metric, or we need to improve the performance of Q-learning and DQN algorithms

### Mean Episode Reward

- Interesting downward trend in DQN, but seems to stabilize, so DQN seems to be converging

DQN did **not** outperform Q-learning and Baseline algorithms, which is consistent with Guss et al. who found DQN only performs better in dense environments.

## Future Work

- Sparse and dense environments for more realistic scenarios and compare DQN
- Incorporate experience replay to improve the performance of DQN
- Hyperparameter tuning to optimize general performance

## References

- Gray, Jonathan, et al. "CraftAssist: A Framework for Dialogue-Enabled Interactive Agents." ArXiv.Org, 2019, arxiv.org/abs/1907.08584. Accessed 24 Oct. 2019.
- Guss, William, et al. NeurIPS 2019 Competition: The MineRL Competition on Sample Efficient Reinforcement Learning Using Human Priors. 2019.
- Szlam, Arthur, et al. Why Build an Assistant in Minecraft? 2019.
- Udagawa, Hiroto, et al. Fighting Zombies in Minecraft With Deep Reinforcement Learning. 2016.