



Solving 2048 Puzzle

Jun Hui Teo ID: junhui96

junhui 96@stanford.edu

Stanford

Motivation

Since games have well-defined utility functions and known rules, using AI to solve games allow for a consistent benchmark to compare inference & learning algorithms.

Challenges

1. Stochastic placement and value of tiles in each turn
2. Identify promising board states efficiently in move generation.

References

Nie, Hou, An, AI Plays 2048: <http://cs229.stanford.edu/proj2016/report/NieHouAn-AIPlays2048-report.pdf>

Yulin Zhou From AlphaGo Zero to 2048 From: <http://www.cse.msu.edu/~zhaoxi35/DRL4KDD/2.pdf>



Stanford University

Problem Definition



Game Dynamics

- Agent chooses one of: UP, DOWN, LEFT, RIGHT
- Tile of value either '2' or '4' placed in an empty tile in each turn
- Merge two tiles of the same value aligned horizontally/vertically to produce a tile of double the value.
- Gain a score of 2^{k+1} if 2 tiles of same value 2^k are merged.

Objective

Choose an action in each turn to:


- Obtain the '2048' tile
- Maximise the game score

Model of problem

Two-player, adversarial, turn-taking, zero sum game with stochastic actions.

Approaches

Depth Limited Expectiminimax Search



$$V_{\text{expectiminimax}}(s, d) =$$

$$\begin{cases} \text{Utility}(s) & \text{IsEnd}(s) \\ \text{Eval}(s) & d = 0 \\ \max_{a \in \text{Action}_{\text{agent}}(s)} V_{\text{expectiminimax}}(\text{Succ}(s, a), d) & \text{Player}(s) = \text{Agent} \\ \min_{a \in \text{Action}_{\text{computer}}(s)} [p * V_{\text{expectiminimax}}(\text{Succ}_{\text{computer}}(s, a, "2"), d-1) + (1-p) * V_{\text{expectiminimax}}(\text{Succ}_{\text{computer}}(s, a, "4"), d-1)] & \text{Player}(s) = \text{Computer} \end{cases}$$

- Learn p via Monte Carlo

$$p = (1 - \eta)p + \eta u, \eta = \frac{1}{\#turns + 1}; u = ['2' \text{ in turn}]$$

- Learn depth d via grid search from d = 1 to 4

Heuristics

4 ¹⁵	4 ¹⁴	4 ¹³	4 ¹²
4 ⁸	4 ⁹	4 ¹⁰	4 ¹¹
4 ⁷	4 ⁶	4 ⁵	4 ⁴
4 ⁰	4 ¹	4 ²	4 ³

S-shaped weight matrix to enforce monotony of aligned tiles

- Pruning of game tree by dynamically ordering successor states by heuristic function.
- Limit branching to a fixed factor f, by selecting top f states for exploration.

Results

For 100 samples, distribution of max tiles are:

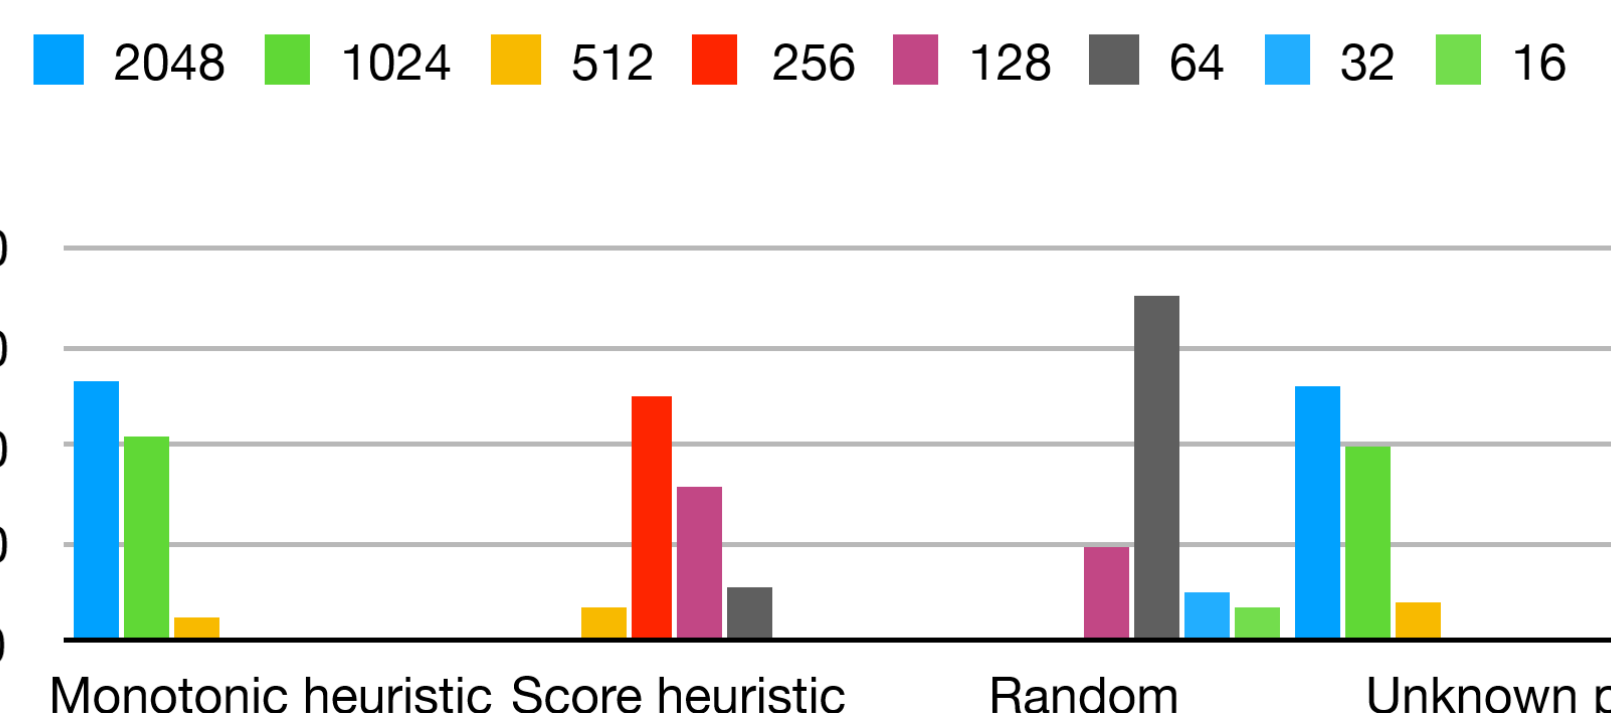


Fig 1: Expectiminimax search to depth 3 with different heuristics v/s baseline

2048 1024 512 256 128




Fig 2: Grid search results with avg normalised time* for each turn

2048 1024 512 256 128




Fig 3: Plot of branching factor f against max tile distribution with normalized time* for a turn

*time is normalised w.r.t d = 1 of unpruned case

Analysis

- Monotonic heuristic performs significantly better than the naive score heuristic or the baseline of random moves - success rate of 53%
- Monte Carlo updates learns p well & has comparable performance.

- Success rate increases with depth as further look-ahead ensures more accurate estimate of minimax value.
- d = 3 best manages the success, time-efficiency trade-off

- Limits to branching early on mean that agent neglects states that seem less promising initially but have higher true minimax values
- Most substantial improvement in performance seen from f=5 to f=6.
- But average time for a step similar to unpruned case, due to time to evaluate heuristic and order states.